

Machine Learning

Dimensionality Reduction

Prof. Matthias Hein

Machine Learning Group
Department of Mathematics and Computer Science
Saarland University, Saarbrücken, Germany


Lecture 26, 05.02.2014

Dimensionality Reduction: Construction of a mapping $\phi : \mathcal{X} \rightarrow \mathbb{R}^m$, where the dimensionality m of the target space is usually much smaller than that of the input space \mathcal{X} . Generally, the mapping should preserve properties of the input space \mathcal{X} e.g. distances.

Why should we do dimensionality reduction ?

- **Manifold assumption:** The internal degrees of freedom are much smaller than the number of measured features \implies data lies along a low-dimensional structure in feature space \implies we want to detect these “true parameters”.
- **Visualization:** interpretation of data in high dimensions is difficult - embeddings in two or three dimensions can provide insight.
- **Data compression:** compress the data but retain most of the information.

Manifold-Assumption



- digits vary smoothly (but discretization as pixels),
- internal degrees of freedom are small compared to the number of features (= number of pixels).

Supervised dimensionality reduction:

- Linear discriminant analysis (LDA),

Unsupervised dimensionality reduction:

- Principal Components Analysis (PCA),
(also called: Karhunen-Loeve-Transformation),
- Kernel PCA,
- Laplacian Eigenmaps,
- Independent Component Analysis (ICA).

Except the last all are eigenvalue problems !

PCA - Two points of view

- the principal k -components span the k -dimensional affine subspace which yields the best approximation of the data (Euclidean norm),
- the subspace spanned by the first k principal components contains “most” of the variance in the data.

PCA - a simple coordinate transformation

- translation - mean of data points becomes new origin,
- rotation - change of the initial ONB into a new ONB which is defined by the data.

PCA - Approximation point of view

Given: $\{X_i\}_{i=1}^n$ in \mathbb{R}^d , Goal: find a m -dimensional affine subspace U_m , with

$$U_m = c + V_m := c + \left\{ \sum_{j=1}^m \alpha_j u_j \mid \{u_j\}_{j=1}^m \text{ ONS}, c \in \mathbb{R}^d, \alpha_j \in \mathbb{R} \right\},$$

which approximates the original data points optimally in the sense,

$$\arg \min_{Z_i \in V_m, c \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \|Z_i + c - X_i\|_2^2.$$

Orthogonal projection P onto the subspace V_m : $P = \sum_{j=1}^m u_j u_j^T$.

Lemma

An orthogonal projection matrix $P : \mathbb{R}^d \rightarrow \mathbb{R}^d$ satisfies,

$$P = P^T, \text{ and } P^2 = P.$$

Optimal offset c

Affine subspace: $U_m = c + V_m$, (c can be seen as origin of U_m).

$$\nabla_c \left(\sum_{i=1}^n \|Z_i + c - X_i\|_2^2 \right) = 2 \sum_{i=1}^n (Z_i - X_i) + 2nc \implies c = \frac{1}{n} \sum_{i=1}^n (X_i - Z_i).$$

- c depends on Z_i - the origin of the subspace U_m can be changed without changing the approximation.
- fix degree of freedom by requiring that

$$\sum_{i=1}^n Z_i = 0 \quad \text{and thus} \quad c = \frac{1}{n} \sum_{i=1}^n X_i.$$

We center the original data points X_i : $\tilde{X}_i = X_i - \frac{1}{n} \sum_{j=1}^n X_j$.

New Objective:
$$\sum_{i=1}^n \|Z_i + c - X_i\|_2^2 = \sum_{i=1}^n \|Z_i - \tilde{X}_i\|_2^2.$$

$$\left\| Z_i - \tilde{X}_i \right\|_2^2 = \left\| Z_i - P\tilde{X}_i \right\|_2^2 + \left\| P\tilde{X}_i - \tilde{X}_i \right\|_2^2,$$

for the orthogonal projection P onto $U_m \implies$ choose $Z_i = P\tilde{X}_i$.

New **transformed objective**:

$$\begin{aligned} \sum_{i=1}^n \left\| Z_i - \tilde{X}_i \right\|_2^2 &= \sum_{i=1}^n \left\| (P - \mathbb{1})\tilde{X}_i \right\|_2^2 \\ &= \sum_{i=1}^n \tilde{X}_i^T (\mathbb{1} - P) \tilde{X}_i \\ &= \sum_{i=1}^n \tilde{X}_i^T \tilde{X}_i - \sum_{i=1}^n \tilde{X}_i^T P \tilde{X}_i \\ &= \sum_{i=1}^n \tilde{X}_i^T \tilde{X}_i - \sum_{j=1}^n u_j^T \left(\sum_{i=1}^n \tilde{X}_i \tilde{X}_i^T \right) u_j \end{aligned}$$

Final objective:



$$\sum_{i=1}^n \|z_i - \tilde{x}_i\|^2 = \sum_{i=1}^n \tilde{x}_i^T \tilde{x}_i - \sum_{j=1}^m u_j^T \left(\sum_{i=1}^n \tilde{x}_i \tilde{x}_i^T \right) u_j.$$

Define the symmetric, positive semi-definite matrix $C \in \mathbb{R}^{d \times d}$ as,

$$C = \sum_{i=1}^n \tilde{x}_i \tilde{x}_i^T,$$

- objective is minimized by using the projection P onto the the m largest eigenvectors of C
- These eigenvectors are called the **principal components** of the data.

PCA - Illustration



- red directions: principal directions in the data
- length of red line: $4\sqrt{\lambda}$, where λ is the eigenvalue of C .

Subspace containing most of the variance of a probability measure

One-dimensional subspace U_1 spanned by $u \Rightarrow$ variance of the data projected onto u is given as

$$\text{var}(u) = \mathbb{E}_X[\langle u, X - \mathbb{E}X \rangle^2] = \mathbb{E}_X \left[(\langle u, X \rangle - \langle u, \mathbb{E}X \rangle)^2 \right].$$

Rewrite $\text{var}(u)$ as

$$\text{var}(u) = \mathbb{E}_X[u^T (X - \mathbb{E}X)(X - \mathbb{E}X)^T u] = \langle u, Cu \rangle,$$

where

$$C = \mathbb{E}_X(X - \mathbb{E}X)(X - \mathbb{E}X)^T,$$

is the **covariance of P_X** .

Subject to $\|u\|^2 = 1 \Rightarrow$ using Rayleigh-Ritz principle, $\text{var}(u)$ is maximized by the eigenvector of C corresponding to the largest eigenvalue.

Best m -dimensional subspace: m “largest” eigenvectors.

- the ev, $\{u_i\}_{i=1}^d$, of C determine an **uncorrelated** ONB,

$$\langle u_i, Cu_j \rangle = \lambda_i \delta_{ij}, \quad i, j = 1, \dots, d.$$

- For **Gaussian** data: $p(x) = \frac{1}{(2\pi)^{\frac{d}{2}} |\det C|^{\frac{1}{2}}} e^{-\frac{1}{2}(x-\mu)^T C^{-1}(x-\mu)}$,
we get in new coordinates z defined as,

$$z = C^{-\frac{1}{2}}(x - \mu) = \sum_{i=1}^d \frac{1}{\sqrt{\lambda_i}} u_i u_i^T (x - \mu),$$

components z_j which are **independent** and equally distributed,

$$p(z) = \frac{1}{(2\pi)^{\frac{d}{2}}} e^{-\frac{\|z\|^2}{2}} = \prod_{j=1}^d \frac{1}{\sqrt{2\pi}} e^{-\frac{z_j^2}{2}}.$$

This process is called **whitening**.

Whitening: PCA + rescaling.

$$z = C^{-\frac{1}{2}}(x - \mu).$$

Whitening are three concatenated operations:

- **centering** - equivalent to a translation in \mathbb{R}^d ,
- **projection onto (all) principal components** - equivalent to a change from the initial basis to the basis spanned by the eigenvectors of C
 \implies rotation,
- **rescaling** - one rescales each axis by the square-root of the corresponding eigenvalue - thus one has unit variance in each direction.

In practice:

- pre-processing of data \implies resulting features are uncorrelated,
- Whitening “spheres” the data - eliminates differences in scaling.

Probability measure unknown only given i.i.d. sample $\{X_i\}_{i=1}^n$
 \implies use **empirical covariance matrix**,

$$C = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T, \quad \text{with} \quad \bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

and use its eigenvalues and eigenvectors as principal components.

Further practical issues:

- never cut the spectrum where two eigenvalues are close,
- several people use the first k -principal components to define new coordinates for supervised problems e.g. classification. This is problematic since the class structure need not have anything to do with the principal components.

Supervised case: use LDA or other supervised extensions of PCA.

Non-linear extension of PCA:

- given: positive definite kernel $k : \mathcal{X} \rightarrow \mathcal{X} \rightarrow \mathbb{R}$,
- map data into the corresponding feature space (RKHS) \mathcal{H}_k ,

$$\phi : \mathcal{X} \rightarrow \mathcal{H}_k, \quad x \rightarrow \phi(x).$$

- do PCA in \mathcal{H}_k (resp. subspace spanned by the data).
- principal components correspond to functions \mathcal{X} .

Questions:

- how to define eigenvectors in \mathcal{H}_k ?
- how many principal components are there ?
- what is a principal component in \mathcal{H}_k ?

Standard-PCA:

$$Cv = \lambda v, \quad \implies \quad \frac{1}{n} \sum_{i=1}^n \langle X_i, v \rangle X_i = \lambda v.$$

\implies all eigenvectors lie in the span of the data points.

Kernel-PCA: map $\phi : \mathcal{X} \rightarrow \mathcal{H}_k$

$$C = \frac{1}{n} \sum_{j=1}^n \phi(X_j) \phi(X_j)^T.$$

If $\dim \mathcal{H}_k = \infty$ then C is a linear operator in \mathcal{H}_k .

As in PCA we want to find the eigenvectors of C ,

$$Cv = \lambda v \quad \implies \quad \frac{1}{n} \sum_{i=1}^n \langle \phi(X_i), v \rangle_{\mathcal{H}_k} \phi(X_i) = \lambda v.$$

\implies all eigenvectors lie in the span of the **mapped** data points.

Kernel PCA - the essential

Kernel-PCA: $Cv = \lambda v \implies \frac{1}{n} \sum_{i=1}^n \langle \phi(X_i), v \rangle_{\mathcal{H}_k} \phi(X_i) = \lambda v.$

Equivalently, solve for all $j = 1, \dots, n$,

$$\frac{1}{n} \sum_{i=1}^n \langle \phi(X_i), v \rangle_{\mathcal{H}_k} \langle \phi(X_i), \phi(X_j) \rangle_{\mathcal{H}_k} = \lambda \langle v, \phi(X_j) \rangle_{\mathcal{H}_k}.$$

Moreover, from the above derivation we know: $v = \sum_{r=1}^n \alpha_r \phi(X_r)$,

$$\frac{1}{n} \sum_{i,r=1}^n \alpha_r \langle \phi(X_i), \phi(X_r) \rangle_{\mathcal{H}_k} \langle \phi(X_i), \phi(X_j) \rangle_{\mathcal{H}_k} = \lambda \sum_{r=1}^n \alpha_r \langle \phi(X_r), \phi(X_j) \rangle_{\mathcal{H}_k}.$$

This can be summarized using $k(X_i, X_j) = \langle \phi(X_i), \phi(X_j) \rangle_{\mathcal{H}_k}$ as,

$$K^T K \alpha = n \lambda K^T \alpha.$$

This is (almost) equivalent to: $K \alpha = n \lambda \alpha.$

What is the difference of the two equations ?

Kernel-PCA: solve eigen-problem: $K\alpha = n\lambda\alpha$.

- normalize eigenvectors $v^{(s)}$, $s = 1, \dots, n$,

$$\left\langle v^{(s)}, v^{(s)} \right\rangle_{\mathcal{H}_k} = \sum_{i,j=1}^n \alpha_i^{(s)} \alpha_j^{(s)} K_{ij} = \lambda^{(s)} \sum_{i=1}^n \alpha_i^{(s)} \alpha_i^{(s)}.$$

- What are the principal components (functions) ? Compute projection of mapped test point x on $v^{(s)}$,


$$\left\langle v^{(s)}, \phi(x) \right\rangle_{\mathcal{H}_k} = \sum_{i=1}^n \alpha_i^{(s)} \langle \phi(X_i), \phi(x) \rangle_{\mathcal{H}_k} = \sum_{i=1}^n \alpha_i^{(s)} k(X_i, x).$$

Standard PCA components are linear functions ! Variation into the direction of the principal component.









- What requirement of PCA did we not integrate into the derivation of Kernel PCA ?

Kernel PCA - Interpretation

Illustration: PCA versus Kernel-PCA




Balanced clusters: Higher principal components of Kernel-PCA




Disbalanced clusters: Higher principal components of Kernel-PCA


Kernel PCA Comp: 1




Kernel PCA Comp: 2




Kernel PCA Comp: 3




Kernel PCA Comp: 4




Kernel PCA Comp: 5




Kernel PCA Comp: 6




Kernel PCA Comp: 7




Kernel PCA Comp: 8



PCA Comp: 1

























PCA Comp: 2



Kernel PCA - Denoising

Kernel-PCA for denoising of data

	Gaussian noise										'speckle' noise									
orig.	0 1 2 3 4 5 6 7 8 9										0 1 2 3 4 5 6 7 8 9									
noisy																				
$n = 1$																				
4																				
16																				
64																				
256																				
$n = 1$																				
4																				
16																				
64																				
256																				

- PCA allows for reconstruction of the original image (just a basis transformation),
- for Kernel PCA this is not directly possible - need to find a pre-image for $\sum_{i=1}^n \alpha_i \phi(x_i) \in \mathcal{H}_k$ in the original space \mathcal{X} .

The continuous Laplacian


$$\mathbb{R}^d, \quad \Delta = \sum_{i=1}^d \frac{\partial^2}{\partial x_i^2}.$$

Why is it interesting ?

- Laplacian is symmetric (self-adjoint),
 - eigenfunctions, $\Delta f = \lambda f$, define an ONB of $L_2(\mathbb{R}^d)$.
 - these eigenfunctions have nice properties
 - ▶ \mathbb{R} : Fourierbasis $\phi_{2k}(x) = \cos(x)$, $\phi_{2k+1}(x) = \sin(x)$,
 - ▶ sphere S^2 : spherical harmonics.
- \implies multi-scale decomposition of the data,
- Fourier-transform is the corresponding basis transformation.

Can we do the same for discrete data ?

The data manifold



- we would like to find the parameters underlying the data-generating process \Rightarrow parameterization of the data-manifold.
- **Idea:** build graph - use graph Laplacian as surrogate of the continuous Laplacian.
 \Rightarrow eigenvectors generate multi-scale decomposition of the data.

Use the graph Laplacian

Three types of graph Laplacians:

unnormalized:
$$(\Delta^{(u)}f)(i) = d(i)f(i) - \sum_{j=1}^n w_{ij}f(j),$$

$$(\Delta^{(u)}f) = (D - W)f,$$

random walk:
$$(\Delta^{(rw)}f)(i) = f(i) - \frac{1}{d(i)} \sum_{j=1}^n w_{ij}f(j),$$

$$(\Delta^{(rw)}f) = (\mathbb{1} - D^{-1}W)f,$$

normalized:
$$(\Delta^{(n)}f)(i) = f(i) - \sum_{j=1}^n \frac{w_{ij}}{\sqrt{d_i d_j}}f(j),$$

$$(\Delta^{(n)}f) = (\mathbb{1} - D^{-1/2}WD^{-1/2})f.$$

Laplacian Eigenmaps

Choose the graph Laplacian: unnormalized, random walk and normalized.


- compute the graph Laplacian $n \times n$ -matrix for n points,
- compute the first k eigenvectors $\{u_i\}_{i=1}^k$ (each eigenvector is normalized, $\|u_i\| = 1$, $i = 1, \dots, k$),
- Embedding $\phi : V \rightarrow \mathbb{R}^k$, of the n vertices into \mathbb{R}^k by $i \rightarrow z_i = (u_1(i), \dots, u_k(i))$,

The embedding: $\phi : V \rightarrow \mathbb{R}^k$, $i \rightarrow \phi(i) = (u_1(i), \dots, u_k(i))$ is the **Laplacian eigenmap**.

Relation to Kernel-PCA:



One can see Laplacian eigenmaps as Kernel PCA with a special data-dependent kernel (pseudo-inverse of the graph Laplacian).

Laplacian Eigenmaps - Computer graphics





- compute eigenvectors of the Laplacian on the mesh,
- can be used for denoising of meshes, varying of meshes etc.

Laplacian Eigenmaps - Illustration



- **Right:** artificial datasets of ones - two variations: line thickness and style variation (bottom line) - digits are of size 28×28 - 784 pixels,
- **Left:** sampling is done uniformly in the parameterization.

Laplacian Eigenmaps - Illustration



- the original parameter set is equivalent to $[0, 1]^2$ and the examples A, B, C, D are the corners of $[0, 1]^2 \implies$ Laplacian eigenmap finds the parameterization.

Independent Component Analysis (ICA)

Motivation: cocktail party problem - blind source separation

- k different speakers (sources),

$$s_1(t), \dots, s_k(t).$$

- d microphones (sensors),

$$x_1(t), \dots, x_d(t).$$

Assumption: measured signal is linear superposition of sources.

Goal: having only the signal of the microphones, find the sources - determine A , where


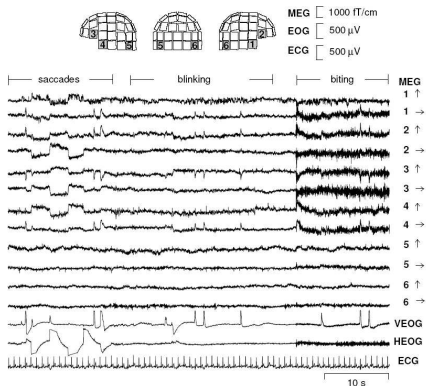
$$x(t) = A s(t).$$

- A is called the **mixing matrix**.


Application scenarios

- sound (speech, music,...) signals,
- EEG signals,
- natural images (patches),
- financial data,
- ...

ICA for EEG analysis



ICA for natural images - 16×16 - patches



- ICA components for 16×16 -patches of natural images,
- \implies one observes that independent components look like edge detectors.

Motivation for ICA

- speakers (sources) are independent of each other.

$$s_1(t), \dots, s_k(t),$$

in the stochastic sense (source signals are independent random variables),

$$p_s(s_1(t), \dots, s_k(t)) = \prod_{i=1}^k p_{s_i}(s_i(t)).$$

Find new representation such that components are maximally independent !

⇒ how can one optimize for independent components ?

⇒ **for simplicity we assume** $d = k$ (nr. sensors = nr. sources).

What kind of independent components can we hope for ?

- **non-Gaussian sources:** suppose that $s(t) \in \mathbb{R}^k$ is Gaussian distributed $\implies x = As$ is again Gaussian distributed,

$$\mathbb{E}[xx^T] = \mathbb{E}[A s s^T A^T] = A \mathbb{E}[s s^T] A^T = A \mathbb{1}_k A^T = A A^T.$$

Whitening yields independent components - but not necessarily $s(t)$.

- **Sources can be identified only up to rescaling:**

$$x(t) = A s(t) = (A D^{-1}) (D s(t)),$$

where D is a diagonal matrix - $D s(t)$ is also independent. W.l.o.g.,

$$\mathbb{E}[s(t) s(t)^T] = \mathbb{1}_k.$$

- **Sources cannot be ordered:** Let P be a permutation matrix, then $P s(t)$ is independent, $x(t) = A s(t) = (A P^{-1}) (P s(t))$.

Whitening as a pre-processing step for ICA

Whitening transforms the signal $x(t)$,

$$y(t) = W x(t) = W A s(t),$$

such it becomes **uncorrelated**,

$$\mathbb{1}_k = \mathbb{E}[y(t) y(t)^T] = \mathbb{E}[W x(t) x(t)^T W^T] = W A \mathbb{E}[s s^T] A^T W^T = W A A^T W^T$$

\implies whitening simplifies the problem since the mixing matrix $W A$ for $y(t)$ is orthogonal.

New problem: find the **orthogonal mixing matrix** $B = W A$

$$y(t) = B s(t),$$

resp. B^T such that $B^T y(t) = B^T B s(t) = s(t)$ is maximally independent.

Steps for ICA:

- apply whitening to the data: $y(t) = W x(t)$.
- find orthogonal de-mixing matrix B s.th. $B y(t)$ is maximally independent.

Different criteria:





- ▶ maximize non-gaussianity of $B y(t)$,
- ▶ minimize mutual information $I(\{B y(t)\}_{i=1}^k)$ - mutual information is zero if and only if joint density of $B y(t)$ factorizes into the product of the marginal densities $\implies B y(t)$ is independent.

Problems:


- joint density of $B y(t)$ hard to estimate \rightarrow problems with mutual inf.
- instead: minimize higher order correlations e.g. **kurtosis**

$$\text{kurt}(y) = \mathbb{E}[y^4] - 3\left(\mathbb{E}[y^2]\right)^2.$$

Illustration of ICA for signal data



ICA - Illustration of ICA



- **Left:** Original sources - individual features are independent $p(x_1, x_2) = p(x_1)p(x_2)$.
- **Middle:** Measured signal - directions of PCA (eigenvectors of covariance matrix) and directions of ICA (columns of estimated mixing matrix) are shown - note that the directions of ICA are **not** orthogonal,
- **Right:** Source signal estimated by ICA - coincides up to rescaling with the original signal.

Cocktail party demo.