Machine Learning
Statistical Learning Theory I

Prof. Matthias Hein

Machine Learning Group
Department of Mathematics and Computer Science
Saarland University, Saarbrücken, Germany

Lecture 27, 06.02.2013

Program of today

A brief overview of results from statistical learning theory

- stochastic convergence,
- different notions of consistency,
- consistency for finite function classes,
- consistency for infinite function classes and the VC dimension,
- universal Bayes consistency - conditions ?
- negative results: no free lunch theorem.

Stochastic Convergence

Motivation

Can we upper bound the deviation of $R(f_n)$ from

- the Bayes risk $R^* = \inf_{f \text{ measurable}} R(f)$
- the best risk $R_{\mathcal{F}} = \inf_{f \in \mathcal{F}} R(f)$ in the class $\mathcal{F}$.

where f_n is the function chosen by the learning algorithm.
Here: Binary classification, canonical zero-one loss.

Concentration

A random variable X is **concentrated** if its distribution is very peaked around the expectation $\mathbb{E}X$ of X.

empirical mean: $\overline{X} = \frac{1}{n} \sum_{i=1}^{n} X_i$, with the $\{X_i\}_{i=1}^{n}$ i.i.d. sample.
Intuition: the distribution of $\overline{X}$ will be concentrated around the true mean $\mathbb{E}\overline{X} = \mathbb{E}X$.

Stochastic Convergence II

Three different notions of convergence of random variables

Definition

Let $\{X_n\}$, $n = 1, 2, \ldots$, be a sequence of random variables. We say that X_n **converges in probability**, $\lim_{n\to\infty} X_n = X \quad$ in probability, if for each $\varepsilon > 0$,

$$\lim_{n\to\infty} \mathrm{P}\big(|X_n - X| \geq \varepsilon\big) = 0.$$

We say that X_n **converges almost surely (with probability 1)**, $\lim_{n\to\infty} X_n = X \quad$ almost surely (a.s.), if

$$\mathrm{P}\big(\omega : \lim_{n\to\infty} X_n(\omega) = X(\omega)\big) = 1.$$

For a fixed $p \geq 1$ we say that X_n **converges in L_p or the p-th mean**, $\lim_{n\to\infty} X_n = X \quad$ in L_p, if

$$\lim_{n\to\infty} \mathbb{E}\big(|X_n - X|^p\big) = 0.$$

Stochastic Convergence III

Proposition

The following implications hold,

- $\lim_{n\to\infty} \mathbb{E}\big(|X_n - X|^p\big) = 0 \quad \Longrightarrow \quad \mathrm{P}\big(|X_n - X| \geq \varepsilon\big) = 0,$
- $\lim_{n\to\infty} X_n = X \quad$ *almost surely* $\quad \Longrightarrow \quad \mathrm{P}\big(|X_n - X| \geq \varepsilon\big) = 0,$
- *If for each* $\varepsilon > 0$,

$$\sum_{n=0}^{\infty} \mathrm{P}\big(|X_n - X| \geq \varepsilon\big) < \infty,$$

then $\lim_{n\to\infty} X_n = X$ *almost surely.*

Relevance for machine learning ?

$R(f_n)$ is a random variable since it depends on the training sample.

- how far is $R(f_n)$ away from the Bayes risk R^* ?
- In which sense $\lim_{n\to\infty} R(f_n) = R^*$?

Consistency (Classification)

Consistency for binary classification:

- Loss function, is 0-1-loss,
- $R(f) = \mathbb{E}\mathbb{1}_{f(X)\neq Y} = \mathrm{P}(f(X) \neq Y)$,
- Bayes risk $R^* = \inf_{f \text{ measurable}} R(f)$.
- best risk in function class $R_{\mathcal{F}} = \inf_{f\in\mathcal{F}} R(f)$ in the class $\mathcal{F}$.

Definition (Consistency)

A classification rule is

- consistent for a distribution of (X, Y) if $\lim_{n\to\infty} R(f_n) = R_{\mathcal{F}}$,
- Bayes consistent for a distribution of (X, Y) if $\lim_{n\to\infty} R(f_n) = R^*$.

We have **weak** (convergence in probability) and **strong** (almost sure convergence) consistency.

The probability $\mathrm{P}\big(R(f_n) - R^* > \varepsilon\big)$ is with respect to all possible training samples of size n.

Consistency (Classification) II

What does consistency mean ?

- The true error of f_n converges to the best possible error,
- asymptotic property - no finite sample statements,
- **distribution dependent**, for example hard margin SVM's are Bayes consistent for distributions where the support of $\mathrm{P}(X|Y=1)$ and $\mathrm{P}(X|Y=-1)$ is linearly separable, but clearly for no problem which is non-separable.

A priori we should make no/too many assumptions about the true nature of the problem !

Universal Consistency

Definition (Universal consistency)

A classification rule/learning algorithm is **universally (weakly/strongly) consistent** if it is (weakly/strongy) consistent for any distribution on $\mathcal{X} \times \mathcal{Y}$.

- strong requirement, since the distribution might be arbitrarily strange.
- nevertheless there exist several universally consistent learning algorithms.

Our main interest: universal consistency

Consistency

Find the best possible function in a class of functions

Every learning algorithm selects either implicitly or explicitly the classifier f_n from some function class $\mathcal{F}$,

Natural decomposition (bias-variance decomposition),

$$R(f_n) - R^* = \underbrace{R(f_n) - \inf_{f \in \mathcal{F}} R(f)}_{\text{Estimation error}} + \underbrace{\inf_{f \in \mathcal{F}} R(f) - R^*}_{\text{Approximation error}} .$$

- The **estimation error** is random since it depends on f_n and thus on the training data - measures the deviation from the best possible risk in the hypothesis class $\mathcal{F}$.
- The **approximation error** is deterministic and measures the deviation of $R_{\mathcal{F}}$ from the Bayes risk R^*. It depends on the hypothesis class $\mathcal{F}$ and the data-generating measure - can only be bounded by making assumptions on the distribution of the data.

Downside of simple function classes

In the worst case we have $R^* = 0$ but $\inf_{f \in \mathcal{F}} R(f) \gg 0$.

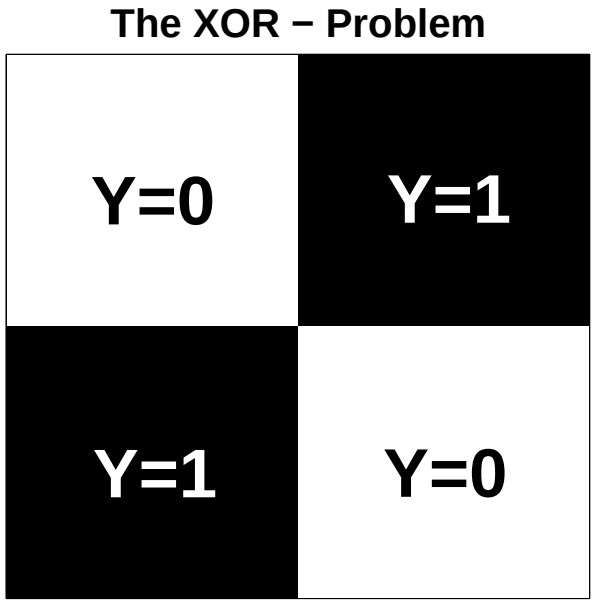

Figure : XOR-problem in $\mathbb{R}^2$. Linear classifiers $\mathcal{F} = \{f(x) = \langle w, x \rangle + b \,|\, w \in \mathbb{R}^2, b \in \mathbb{R}\}$ are very bad but $R^* = 0$.

The basic principle

Proposition

Let f_n be chosen by empirical risk minimization, that is $f_n = \arg\min_{f \in \mathcal{F}} R_n(f)$ where $R_n(f) = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{f(X_i) \neq Y_i}$. Then

$$R(f_n) - \inf_{f \in F} R(f) \leq 2 \sup_{f \in F} |R(f) - R_n(f)|.$$

Proof: We have with $f_{\mathcal{F}}^* = \arg\min_{f \in \mathcal{F}} R(f)$,

$$\begin{aligned} R(f_n) - \inf_{f \in \mathcal{F}} R(f) &= R(f_n) - R_n(f_n) + R_n(f_n) - R(f_{\mathcal{F}}^*) \\ &\leq R(f_n) - R_n(f_n) + R_n(f_{\mathcal{F}}^*) - R(f_{\mathcal{F}}^*) \\ &\leq 2 \sup_{f \in \mathcal{F}} |R_n(f) - R(f)|, \end{aligned}$$

where the second inequality follows from the fact that f_n minimizes the empirical risk.

Definition of empirical processes

Definition

A **stochastic process** is a collection of random variables $\{Z_n, n \in T\}$ on the same probability space, indexed by an arbitrary index set T. An **empirical process** is a stochastic process based on a random sample.

In statistical learning theory we are studying the empirical process,

$$\sup_{f \in \mathcal{F}} |R_n(f) - R(f)|,$$

since uniform control of the deviation $R_n(f) - R(f)$ yields consistency !

$$R(f_n) - \inf_{f \in F} R(f) \leq 2 \sup_{f \in F} |R(f) - R_n(f)|.$$

Hoeffding's inequality

Theorem

Let $X_1, \ldots, X_n$ be independent, bounded and identically distributed random variables such that X_i falls in the interval $[a_i, b_i]$ with probability one. Then for any $\varepsilon > 0$ we have

$$\mathrm{P}\Big(\Big|\frac{1}{n}\sum_{i=1}^{n} X_i - \frac{1}{n}\sum_{i=1}^{n} \mathbb{E}X_i\Big| \geq \varepsilon\Big) \leq 2\exp\Big(-\frac{2n\varepsilon^2}{\frac{1}{n}\sum_{i=1}^{n}(b_i - a_i)^2}\Big).$$

Control of the deviation for a **fixed** function with $R(f) = \mathbb{E}[\mathbb{1}_{f(X)\neq Y}]$,

$$\mathrm{P}\Big(\Big|R_n(f) - R(f)\Big| \geq \varepsilon\Big) \leq 2\exp\Big(-2n\varepsilon^2\Big).$$

Important: This cannot be simply applied to f_n - the function found by empirical risk minimization - since f_n depends on the training data.

A finite set of functions

Bounds for the case of a finite set of functions $\mathcal{F}$

Proposition

Let $\mathcal{F}$ be a finite set of functions, then

$$\mathrm{P}\Big(\sup_{f \in \mathcal{F}} \Big| R_n(f) - R(f) \Big| \geq \varepsilon \Big) \leq 2|\mathcal{F}| \exp\Big(-2n\varepsilon^2 \Big),$$

where $|\mathcal{F}|$ is the cardinality of $\mathcal{F}$. And thus with probability $1-\delta$,

$$R(f_n) \leq R(f_{\mathcal{F}}^*) + \sqrt{\frac{\log|\mathcal{F}| + \log\frac{2}{\delta}}{n}}.$$

Proof: Noting that $0 \leq \mathbb{1}_{f(X)\neq Y} \leq 1$ we get the result using Hoeffding's inequality. Then with $\delta = 2|\mathcal{F}|e^{-2n\varepsilon^2}$ one gets $\varepsilon = \sqrt{\frac{1}{n}\Big(\log|\mathcal{F}| + \log\frac{2}{\delta} \Big)}$.

The **convergence rate** is of order $\frac{1}{\sqrt{n}} \Longrightarrow$ typical in SLT.

Infinite number of functions

Major contribution of Vapnik and Chervonenkis: uniform deviation bounds over general infinite classes.

Given points $x_1, \ldots, x_n$ and a class $\mathcal{F}$ of binary-valued functions denote by

$$\mathcal{F}_{x_1,\ldots,x_n} = \Big\{ \{f(x_1), \ldots, f(x_n)\} \,|\, f \in \mathcal{F} \Big\},$$

the set of all possible classification of the set of points via functions in $\mathcal{F}$.

Definition

The **growth function** $S_{\mathcal{F}}(n)$ is the maximum number of ways into which n points can be classified by the function class $\mathcal{F}$,

$$S_{\mathcal{F}}(n) = \sup_{(x_1,\ldots,x_n)} |\mathcal{F}_{x_1,\ldots,x_n}|.$$

If $S_{\mathcal{F}}(n) = 2^n$ we say that $\mathcal{F}$ **shatters** n points.

Why is this growth function interesting ?

Symmetrization lemma

- **ghost sample:** a second i.i.d. sample of size n (independent of the training data).
- $R'_n(f)$ denotes the empirical risk associated with the ghost sample.

Lemma

Let $n\varepsilon^2 \geq 2$, *we have*

$$\mathrm{P}\left(\sup_{f \in \mathcal{F}} |R_n(f) - R(f)| > \varepsilon \right) \leq 2\,\mathrm{P}\left(\sup_{f \in \mathcal{F}} |R_n(f) - R'_n(f)| > \frac{\varepsilon}{2} \right),$$

- **Important:** $|R_n(f) - R'_n(f)|$ depends only on the values of the function takes on the $2n$ samples - these are maximum 2^{2n} different values $\Longrightarrow$ independent of how many functions are contained in $\mathcal{F}$.
- a simple union bound will now yield the V(apnik)C(hervonenkis)-bound.

VC Bound for general $\mathcal{F}$

The growth function is a measure of the "size" of $\mathcal{F}$,

Theorem (Vapnik-Chervonenkis)

For any $\delta > 0$, with probability at least $1 - \delta$,

$$\forall f \in \mathcal{F}, \qquad R(f_n) \leq R(f_{\mathcal{F}}^*) + 8\sqrt{\frac{\log S_{\mathcal{F}}(2n) + \log \frac{8}{\delta}}{2n}}.$$

Proof:

$$\begin{aligned}
\mathrm{P}\big(R(f_n) - \inf_{f\in\mathcal{F}} R(f) > \varepsilon\big) &\leq \mathrm{P}\big(\sup_{f\in F} |R(f) - R_n(f)| > \frac{\varepsilon}{2}\big) \\
&\leq 2\,\mathrm{P}\Big(\sup_{f\in\mathcal{F}} |R_n(f) - R_n'(f)| > \frac{\varepsilon}{4}\Big) \\
&\leq 2\,S_{\mathcal{F}}(2n)\,\mathrm{P}\Big(|R_n(f) - R_n'(f)| > \frac{\varepsilon}{4}\Big) \\
&\leq 4\,S_{\mathcal{F}}(2n)\,\mathrm{P}\Big(|R_n(f) - R(f)| > \frac{\varepsilon}{8}\Big) \leq 8\,S_{\mathcal{F}}(2n)\,e^{-\frac{n\varepsilon^2}{32}}
\end{aligned}$$

Discussion of VC-Bound

For a finite class $\log S_{\mathcal{F}}(n) \leq |\mathcal{F}| \Rightarrow$ up to constants at least as good as the previous bound for finite $\mathcal{F}$.

Definition

The **VC dimension** $\mathrm{VC}(\mathcal{F})$ of a class $\mathcal{F}$ is the largest n such that $S_{\mathcal{F}}(n) = 2^n$.

What happens if $\mathcal{F}$ can always realize all 2^n possibilities ?

$$R(f_n) \leq R(f_{\mathcal{F}}^*) + 8\sqrt{\frac{\log S_{\mathcal{F}}(2n) + \log \frac{8}{\delta}}{2n}}$$

$$\leq R(f_{\mathcal{F}}^*) + 8\sqrt{\frac{n \log 2 + \log \frac{8}{\delta}}{2n}}$$

The second term does not converge to zero as $n \to \infty$!

$\Longrightarrow$ bound suggests that restricted $\mathcal{F}$ is required for generalization.

VC dimension

What happens with $S_{\mathcal{F}}(n)$ for $n > \mathrm{VC}(\mathcal{F})$?

We know: $n \leq \mathrm{VC}(\mathcal{F}) \implies S_{\mathcal{F}}(n) = 2^n$ but what if $n > \mathrm{VC}(\mathcal{F})$?

Lemma (Vapnik-Chervonenkis, Sauer, Shelah)

Let $\mathcal{F}$ be a class of functions with finite VC-dimension $\mathrm{VC}(\mathcal{F})$. Then for all $n \in \mathbb{N}$,

$$S_{\mathcal{F}}(n) \leq \sum_{i=0}^{\mathrm{VC}(\mathcal{F})} \binom{n}{i},$$

and for all $n > \mathrm{VC}(\mathcal{F})$,

$$S_{\mathcal{F}}(n) \leq \left(\frac{e\, n}{\mathrm{VC}(\mathcal{F})}\right)^{\mathrm{VC}(\mathcal{F})}.$$

Phase transition from exponential to polynomial growth of $S_{\mathcal{F}}(n)$

VC bound II

Plugging the bounds on the growth function into the VC bounds

Corollary

Let $\mathcal{F}$ be a function class with VC-dimension $\mathrm{VC}(\mathcal{F})$, then for $2n > \mathrm{VC}(\mathcal{F})$ one has for any $\delta > 0$, with probability at least $1 - \delta$,

$$\forall f \in \mathcal{F}, \qquad R(f_n) \leq R(f_{\mathcal{F}}^*) + 8\sqrt{\frac{\mathrm{VC}(\mathcal{F}) \log \frac{2en}{\mathrm{VC}(\mathcal{F})} + \log \frac{8}{\delta}}{2n}}.$$

Deviation of $R(f_n)$ from $R(f_{\mathcal{F}}^*) = \inf_{f \in \mathcal{F}} R(f)$ decays as $\sqrt{\mathrm{VC}(\mathcal{F}) \frac{\log n}{n}}$.

- VC dimension is not just counting the number of functions but the variability of the functions in the class on the sample.
- finite VC dimension ensures **universal consistency**,
- other techniques for bounds exist: covering numbers, Rademacher averages.

VC bound III

Necessary and sufficient conditions for consistency

The following theorem is one of the key-theorems for statistical learning.

Theorem (Vapnik-Chervonenkis (1971))

A **necessary** *and* **sufficient** *condition for the universal consistency of empirical risk minimization using a function class $\mathcal{F}$ is,*

$$\lim_{n\to\infty} \frac{\log S_{\mathcal{F}}(n)}{n} = 0.$$

We have proven that $\lim_{n\to\infty} \frac{\log S_{\mathcal{F}}(n)}{n} = 0$ is sufficient for consistency.
The proof, that this condition is also necessary requires a bit more effort.

Is the restriction necessary ?

Empirical risk minimization can be inconsistent

Input space: $\mathcal{X} = [0, 1]$. The labels are deterministic

$$Y = \begin{cases} -1, & \text{if } X \leq 0.5, \\ 1, & \text{if } X > 0.5. \end{cases}$$

We consider the following classifier,

$$f_n(X) = \begin{cases} Y_i & \text{if } X = X_i \text{ for some } i = 1, \ldots, n \\ 1 & \text{otherwise.} \end{cases} .$$

We have $R_n(f_n) = 0 \qquad \text{but} \qquad R(f_n) = \frac{1}{2}$.
The classifier f_n is **not Bayes consistent**. We have,

$$\lim_{n \to \infty} R(f_n) = \frac{1}{2} \neq 0 = R^*.$$

$\Longrightarrow$ just memorizing - no learning, no generalization.

VC Dimension

VC dimensions of selected function classes:

- The set of linear halfspaces in $\mathbb{R}^d$ has VC dimension $d + 1$.
- The set of linear halfspaces of margin ρ and where the smallest sphere enclosing the data has radius R has VC dimension,

$$\mathrm{VC}(\mathcal{F}) \leq \min\left\{d, \frac{4\,R^2}{\rho^2}\right\} + 1.$$

- The function $\mathrm{sign}(\sin(tx))$ on $\mathbb{R}$ has infinite VC dimension.

$\Rightarrow$ VC dimension has nothing to do with the number of free parameters !

Justification for Support Vector machines

The set of linear halfspaces of margin ρ and where the smallest sphere enclosing the data has radius R has VC dimension,

$$\mathrm{VC}(\mathcal{F}) \leq \min\left\{d, \frac{4\,R^2}{\rho^2}\right\} + 1.$$

The vector w of the optimal maximal-margin hyperplane satisfies,

$$\|w\|^2 = \frac{1}{\rho^2},$$

Thus, the Support-Vector Machine (SVM)

$$\min_{w,b} \frac{1}{n} \sum_{i=1}^{n} \max\{0, 1 - Y_i(\langle w, X_i\rangle + b)\} + \lambda \|w\|^2 .$$

penalizes large margins $\|w\| \Longrightarrow$ **limits capacity of function class**

Remarks on VC bounds (applies also to other existing bounds)

- **No a-posteriori justification:** bounds cannot be used for a posteriori justification. In particular, the bound holds not for the margin obtained by the SVM, but the bound holds for a function class with pre-defined margin (before seeing the data) !
- **Bounds are often loose:** the bounds are **worst-case bounds** which apply to **any** possible probability measure on $\mathcal{X} \times \mathcal{Y} \Longrightarrow$ for practical sample sizes bounds are often larger than 1 ! But: bounds capture certain characteristics of the learning algorithm.

Universal Bayes consistency

Decomposition into estimation and approximation error),

$$R(f_n) - R^* = \underbrace{R(f_n) - \inf_{f \in \mathcal{F}} R(f)}_{\text{Estimation error}} + \underbrace{\inf_{f \in \mathcal{F}} R(f) - R^*}_{\text{Approximation error}} .$$

$\Longrightarrow$ up to now fixed function class $\Longrightarrow$ fixed approximation error.

Structural risk minimization:

- Let the function class $\mathcal{F}$ be a function of the sample size n: $\mathcal{F}_n$.
- as $n \to \infty$ let $\mathcal{F}_n$ grow so that in the limit it can model any function but estimation error is still bounded:

$$\forall\, f \in \mathcal{F}_n, \qquad R(f_n) \leq R(f_{\mathcal{F}}^*) + 8\sqrt{\frac{\mathrm{VC}(\mathcal{F}_n) \log \frac{2\,e\,n}{\mathrm{VC}(\mathcal{F}_n)} + \log \frac{8}{\delta}}{2n}}.$$

$\Longrightarrow$ **Universal Bayes consistency**

Questions

Naturally arising questions

- Can we quantify the convergence to the Bayes risk ? Can we obtain rates of convergence ?
- What does universal consistency mean for the finite sample case ?
- Is there a universally best learning algorithm ?

No free lunch I

First negative result

Intuition: For every fixed n there exists a distribution where the classifier is arbitrarily bad !

Theorem

For any $\varepsilon > 0$ and any integer n and classification rule f_n, there exists a distribution of (X, Y) with Bayes risk $R^* = 0$ such that

$$\mathbb{E}R(f_n) \geq \frac{1}{2} - \varepsilon.$$

- construct a distribution on the set $\mathcal{X} = \{1, \ldots, K\}$,
- noise-free but no structure,
- for fixed n choose K sufficiently large such that the rule f_n will fail completely on the rest of $\mathcal{X}$.

No free lunch II

First negative result

There exists no universally consistent learning algorithm such that $R(f_n)$ converges uniformly over all distributions to R^*.

No free lunch III

Second negative result

Theorem

Let $\{a_n\}$ be a sequence of positive numbers converging to zero with $\frac{1}{16} \geq a_1 \geq a_2 \geq \ldots$. For every sequence of classification rules, there exists a distribution of (X, Y) with $R^* = 0$, such that for all n,

$$\mathbb{E}R(f_n) \geq a_n.$$

This result states that universally good learning algorithms do not exist
$\Rightarrow$ convergence to the Bayes risk can be **arbitrarily slow** !

There exist no universal rates to the Bayes risk. If one wants to have rates of convergence to the Bayes risk one has to restrict the class of distributions on $\mathcal{X} \times \mathcal{Y}$.

No free lunch IV

Third negative result

Theorem

For every sequence of classification rules f_n, there is a universally consistent sequence of classification rules g_n such that for some distribution on $\mathcal{X} \times \mathcal{Y}$

$$\mathrm{P}\big(f_n(X) \neq Y\big) > \mathrm{P}\big(g_n(X) \neq Y\big), \quad \forall n \geq 0.$$

Thus for every universally consistent learning rule there exists a distribution on $\mathcal{X} \times \mathcal{Y}$ such that another universally consistent learning rule is strictly better.

There exists no universally superior learning algorithm.

No free lunch V

Summary

1. Restriction of the class of distributions on $\mathcal{X} \times \mathcal{Y}$ $\Longrightarrow$ convergence rates to Bayes for universally consistent learning algorithms.
   Problem: Assumptions cannot be tested. Performance guarantees are only valid under the made assumptions.
2. Restriction of the function class $\Longrightarrow$ no universal consistency possible.
   Comparison to the best possible function in the class is possible uniformly over all distributions.
   But **no performance guarantees** with respect to the Bayes risk.

Convergence rates to Bayes

Convergence rates to Bayes only possible under assumptions on the distribution of (X, Y)

Reasonable assumptions fulfill two requirements:

- The assumptions should be as natural as possible, meaning that one expects that most the data generating distributions one encounters in nature fulfill these assumptions.
- The assumptions should be narrow enough, so that one can still prove convergence rates.

Convergence rates to Bayes II

Assumptions

In terms of the regression function: $\eta(x) = \mathbb{E}[Y|X = x]$.

- $\eta(x)$ lies in some Sobolev space (has certain smoothness properties),
- Margin/low noise conditions introduced by Massart and Tsybakov,

Definition

A distribution P on $\mathcal{X} \times \{-1, 1\}$ fulfills the low noise condition if there exist constants $C > 0$ and $\alpha \geq 0$ such that

$$\mathrm{P}\big(|\eta(X)| \leq t\big) \leq Ct^{\alpha}, \qquad \forall\, t \geq 0.$$

The coefficient α is called the **noise coefficient** of P.

1. $\alpha = 0$ is trival and implies no restrictions on the distribution,
2. $\alpha = \infty$, $\eta(x)$ strictly bounded away from zero.

Universal consistency for soft-margin SVM's

Definition

A continuous kernel $k : \mathcal{X} \times \mathcal{X} \to \mathbb{R}$ is called **universal** if the associated RKHS $\mathcal{H}_k$ is dense in the set of continuous functions $C(X)$ with the $\|\cdot\|_\infty$-norm, that is for all $f \in C(X)$ and $\varepsilon > 0$ there exists a $g \in \mathcal{H}_k$ such that

$$\|f - g\|_\infty \leq \varepsilon.$$

$\Rightarrow$ Measurable functions can be approximated by continuous functions.

A soft-margin SVM in $\mathbb{R}^d$ with a **universal kernel** is universally consistent.

Theorem

Let $\mathcal{X} \subset \mathbb{R}^d$ be compact, then the soft-margin SVM with error parameter $C_n = n^{1-\beta}$ for some $0 < \beta < \frac{1}{d}$ and a Gaussian kernel is universally

Large scale empirical risk minimization

Integrate accuracy of optimization (Bottou and Bousquet(2011))

- instead of empirical risk minimizer f_n compute approximation $\tilde{f}_n$

$$R_n(\tilde{f}_n) \leq R_n(f_n) + \rho,$$

where ρ is the tolerance/accuracy with which we do the optmization

New decomposition of the excess risk:

$$R(\tilde{f}_n) - R^* = \underbrace{\inf_{f \in \mathcal{F}} R(f) - R^*}_{\text{Approximation error}} + \underbrace{R(f_n) - \inf_{f \in \mathcal{F}} R(f)}_{\text{Estimation error}} + \underbrace{R(\tilde{f}_n) - R(f_n)}_{\text{Optimization error}}.$$

Learning under fixed budget ($n < n_{\max}$, time $T < T_{\max}$)

- *Small-scale Learning (restricted by $n_{\max}$): as $n_{\max}$ is small the computing time $T(\mathcal{F}, n, \rho)$ is small and thus one can afford ρ to be small (standard setting)*
- *Large-scale Learning (restricted by $T_{\max}$): $T(\mathcal{F}, n, \rho) < T_{\max}$ can be achieved either for ρ small only if $n \ll n_{\max}$ is small (high estimation error). Instead use larger ρ (higher optimization error) and inspect more data points n (lower estimation error)*

Gradient descent versus stochastic gradient descent

Goal: minimize empirical loss: $\frac{1}{n}\sum_{i=1}^{n} L(f_w(x_i), y_i)$

- **Gradient Descent:** At each iteration update w as

$$w_{t+1} = w_t - \frac{\eta}{n}\sum_{i=1}^{n} \nabla_w L(f_w(x_i), y_i).$$

- **Stochastic Gradient Descent:** At each iteration t pick a random training example (x_t, y_t) and update w as

$$w_{t+1} = w_t - \frac{\eta}{t}\nabla_w L(f_w(x_t), y_t).$$

Alg.	Cost per Iteration	Iterations to reach ρ	Time to reach ρ	Time to reach excess risk of ε
GD	$O(nd)$	$O\left(\kappa \log\left(\frac{1}{\rho}\right)\right)$	$O\left(nd\,\kappa \log\left(\frac{1}{\rho}\right)\right)$	$O\left(\frac{d^2\kappa}{\varepsilon^{\frac{1}{\alpha}}} \log\left(\frac{1}{\rho}\right)\right)$
SGD	$O(d)$	$\frac{\nu\kappa^2}{\rho} + o\left(\log\left(\frac{1}{\rho}\right)\right)$	$O\left(\frac{d\nu\kappa}{\rho^2}\right)$	$O\left(\frac{d\kappa^2\nu}{\varepsilon}\right)$

κ: condition number of Hessian, ν : related complexity parameter, α : decay rate of estimation err $(\frac{1}{2} < \alpha < 1)$

Experiment

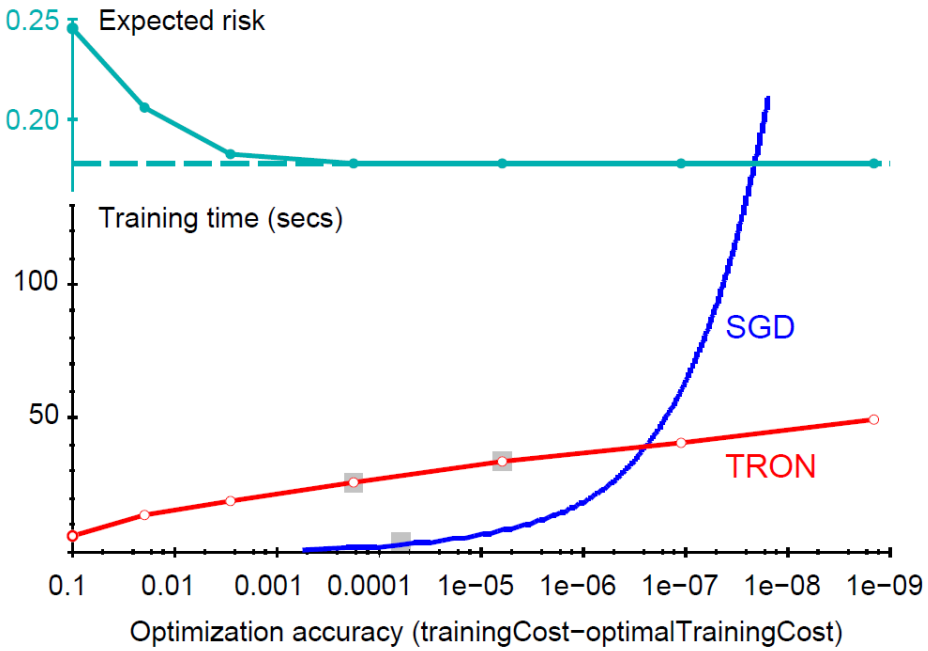
0.25
Expected risk
0.20
Training time (secs)
100
50
SGD
TRON
0.1
0.01
0.001
0.0001
1e−05
1e−06
1e−07
1e−08
1e−09
Optimization accuracy (trainingCost−optimalTrainingCost)

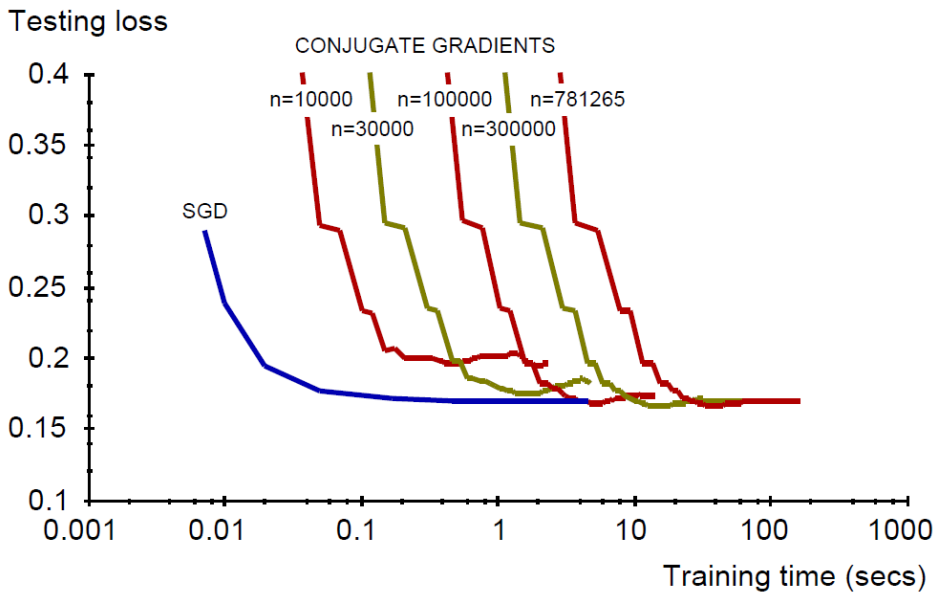
Testing loss
CONJUGATE GRADIENTS
n=10000
n=100000
n=781265
n=30000
n=300000
SGD
0.4
0.35
0.3
0.25
0.2
0.15
0.1
0.001
0.01
0.1
1
10
100
1000
Training time (secs)

Bachelor/Master/PhD topics in machine learning !

Thanks for your attention !